# Final Homework Report

NPFL054 Introduction to Machine Learning

Milan Wikarski

Task 1 – Data analysis

Dataset consists of 5822 examples. Each example has 86 attributes. The 86-th attribute *Purchase* with values *Yes (1)* and *No (0)* is the targte attribute.

Target feature distribution

The distribution of the target attribute *Purchase* is heavily skewed towards *No*. The frequency of *Yes* is just 5.98%. This means that the expected precision when randomly selecting 100 examples whould be 5.98%.

Task 1a

MOSHOODF

Attribute *MOSHOODF* represents the main customer type. This attribute divides customers into 10 groups described in *L2*. The table below lists the number of customers per group, as well as the percentage of examples in this group with target attribute *Purchase* value *Yes (1)* (ie. the percentage of customers who have purchased the caravan insurance policy):

ID	Group	Size	Purchase Frequency
1	Successfull Hedonists	552	8.7%
2	Driven Growers	502	13.15%
3	Average Family	886	6.66%
4	Career Loners	52	0%
5	Living Well	569	2.64%
6	Cruising Seniors	205	1.95%
7	Retired and Religious	550	3.64%

ID	Group	Size	Purchase Frequency
8	Family with Grown ups	1563	5.69%
9	Convervative Families	667	6.3%
10	Farmers	276	1.81%

This data can also be plotted into two barcharts:

Customer main type distribution

Customer main type ~ Purchase

We can further analyze irregularities in groups by plotting the data into two boxplots and looking for outliers:

Group sizes boxplot

Purchase frequency per group boxplot

Group *4 (Career Loners)* is really small (only 52 members) and contains no positive examples but cannot be classified as outlier. Group *8 Family with grown ups* is the largest (1563 members), thus can be classified as outlier base on size, but the *Purchase* frequency of this group is slightly below average (5.69%). When it comes to *Purchase* frequncy in groups, there are no outliers.

MOSTYPE

Attribute *MOSTYPE* represents the customer subtype. This attribute divides customers into 41 subgroups described in *L0*. The table below lists the number of customers per subgroup, as well as the percentage of examples in this subgroup with target attribute *Purchase* value *Yes (1)*:

ID	Subgroup	Size	Purchase Frequency
1	High Income	124	10.48%
2	Very Important Provincials	82	7.32%
3	High status seniors	249	10.04%
4	Affluent senior apartments	52	3.85%
5	Mixed seniors	45	4.44%
6	Career and childcare	119	10.08%
7	Dinki's (double income no kids)	44	6.82%
8	Middle class families	339	15.04%
9	Modern	278	4.32%
10	Stable family	165	5.45%
11	Family starters	153	5.88%
12	Affluent young families	111	14.41%
13	Young all american family	179	7.26%
15	Senior cosmopolitans	5	0%
16	Students in apartments	16	0%
17	Fresh masters in the city	9	0%
18	Single youth	19	0%
19	Suburban youth	3	0%
20	Etnically diverse	25	8%
21	Young urban have-nots	15	0%
22	Mixed apartment dwellers	98	4.08%
23	Young and rising	251	1.59%
24	Young	180	2.78%
25	Young seniors in the city	82	2.44%
26	Own home elderly	48	2.08%
27	Seniors in apartments	50	2%
28	Residential elderly	25	0%
29	Porchless seniors: no front yard	86	2.33%
30	Religious elderly singles	118	3.39%
31	Low income catholics	205	2.93%
32	Mixed seniors	141	5.67%
33	Lower class large families	810	5.68%
34	Large family	182	4.95%
35	Village families	214	3.74%
36	Couples with teens 'Married with children'	225	7.11%
37	Mixed small town dwellers	132	7.58%
38	Traditional families	339	6.78%
39	Large religous families	328	5.79%
40	Large family farms	71	0%
41	Mixed rurals	205	2.44%

The subgroup sizes can also be visualized using a barchart:

Customer subtype distribution

The percentage of people in each subgroup who have purchased the caravan insurance policy can also be plotted to a barchart:

Customer subtype ~ Purchase

It should be noted that among 5822 examples, there is no representative of subgroup *14 (Junior cosmopolitan)*. To further see, which subgroups might be interessting, we can plot boxplots for size and *Purchase* frequency:

Subgroup sizes boxplot

Purchase frequency per group boxplot
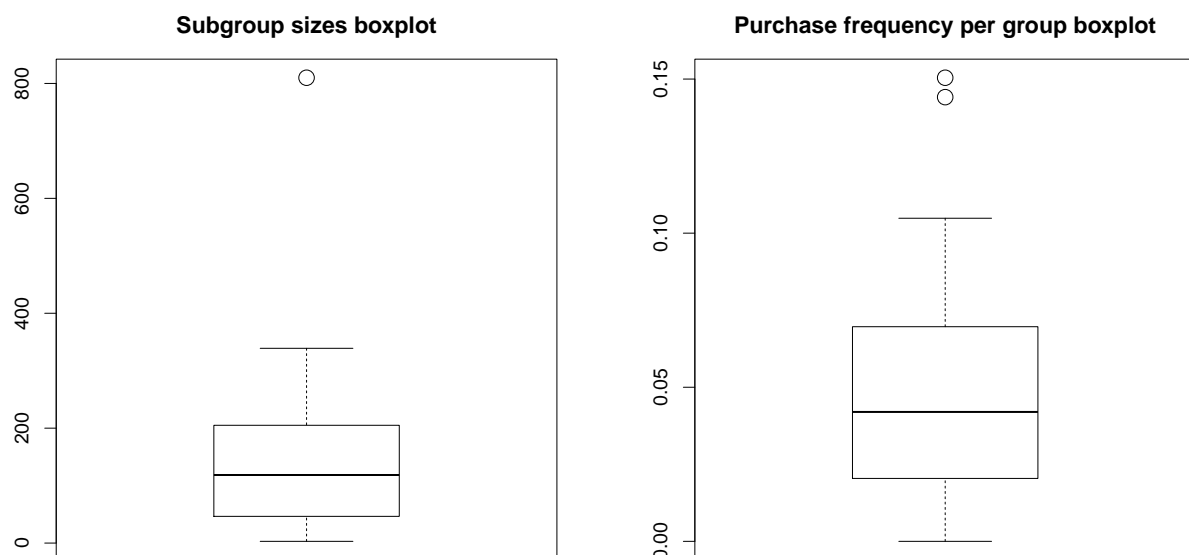
Subgroups *15 − 19* are very small and contain no positive examples but cannot be classified as outliers (the reason for this will be explained in the next chapter). Other subgroups with no positive examples are subgroup *21 Young urban have*-nots, subgroup *28 Residential elderly* and subgroup

40 Large family farms. Subgroups *21* and *28* are too small to draw any conclusions but there seems to be a correlation between being a member of subgroup *40* and not purchasing the caravan insurance policy.

The largest subgroup *33 Lower class large families* is composed of almost 14% of all examples but the *Purchase* frequency in this subgroup is slightly below the average (5.68%). This subgroup is an outlier based on size.

The most promising subgroups are subgroup *12 Affluent young families* with 111 members and 14.41% Purchase frequency, and subgroup *8 Middle class families* with 339 members and *Purchase* frequency 15.04%, which is the greatest among all subgroups. These two subgroups can be classified as outliers based on *Purchase* frequency.

Task 1b

After some analysis, it can be seen that *MOSHOODF* divides the customers into 10 groups and *MOSTYPE* further divides these customers into subgroups. This information can be gained by calling:

```
table(MOSTYPE, MOSHOOFD)
```

and analyzing the output:

```
        MOSHOOFD
MOSTYPE   1    2    3    4    5    6    7    8    9   10
     1  124    0    0    0    0    0    0    0    0    0
     2   82    0    0    0    0    0    0    0    0    0
     3  249    0    0    0    0    0    0    0    0    0
     4   52    0    0    0    0    0    0    0    0    0
     5   45    0    0    0    0    0    0    0    0    0
     6    0  119    0    0    0    0    0    0    0    0
     7    0   44    0    0    0    0    0    0    0    0
     8    0  339    0    0    0    0    0    0    0    0
     9    0    0  278    0    0    0    0    0    0    0
    10    0    0  165    0    0    0    0    0    0    0
   ...
```
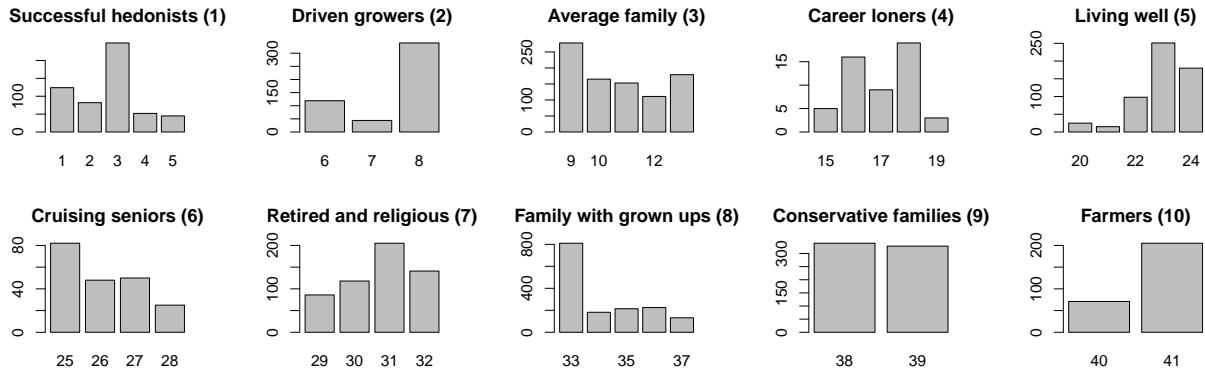
It is clear that all customers of certain subgroup *(MOSTYPE)* belong to just one group (*MOSHOOFD)* and that every customer from each group (*MOSHOOFD)* is assigned one subgroup (*MOSTYPE)*.
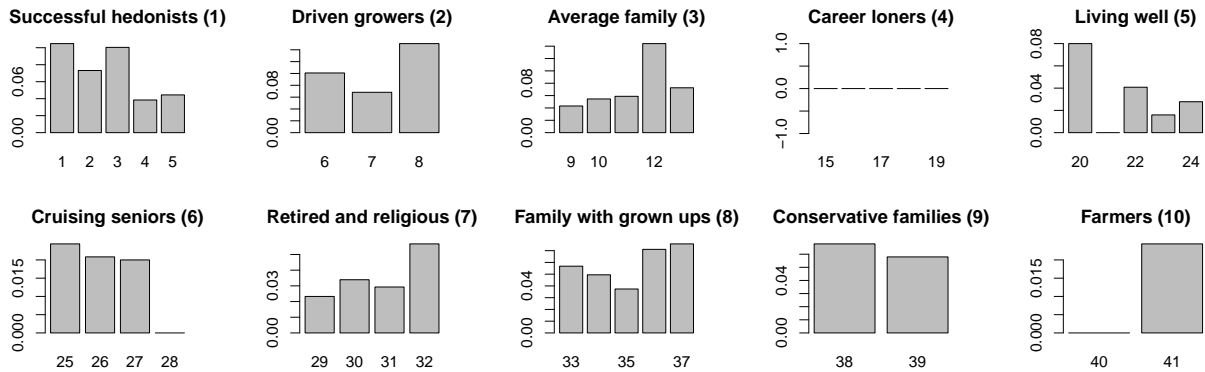
With this information, we can explore the groups (*MOSHOOFD)* in more detail by looking at the size of its subgroups:

Distribution of subgroups in each of customer main type groups

and the *Purchase* frequency of its subgroups:

Purchase frequency of subgroups in each of customer main type groups

This visuallization is really useful when exlaining why subgroups *15-19* contain no examples. These subgroups are part of group *4 Career loners,* which as was shown in **Task 1a** contains no positive examples.

Groups which are interesting are the ones with irregularities in the *Purchase* distribution. Most notably, the likelihood of the members of group *3 Average family* to purchase the caravan insurance policy is average (6.66%), unless they are part of subgroup *12 Affluent young families,* in which case it is 14.41%.

Average family (3)

Subgroups sizes

Purchase frequency

The same can be said about members of group *5 Living well* and subgroup *20 Ethnically diverse*. Although this subgroup can be classified as an outlier compared to other subgrous in the parent group (based on *Purchase* frequency), it contains only 25 members which drastically decreases its importance.

Living well (5)

Subgroups sizes

Purchase frequency

Charts like these two (group details) were created for each group and can be seen in */out/group-<group number>-detail.pdf*.

Task 2 – Model fitting, optimization and selection

Three models were fitted, optimized and compared: *Decision Tree (2a)*, *Random Forest (2b)* and *Regularized Logistic Regression (2c)*.

Each model was fitted and optimized using a similar method: model performance with certain parameters was measured by computing the mean of $AUC_{0.2}$ of 10-fold cross-validation's ROC. Parameters which yielded the best performance were selected, taking confidence intervals into consideration.

Afterwads, models with best parameters were evaluated using the train data set. The model which produced the best results (again, evaluated using $AUC_{0.2}$) was selected to be the final model.

Cross-validation

It is important to note that cross-validation was regularized by ensuring that same number of positive examples was in each fold. For this purpuse a custom function *cv.split.safe()* was defined.

Parameter tuning data

Parameters for a certain model are selected by creating a large number of models, one for each value from a predefined range and then running a 10-fold cross-validation (thus creating and testing 10 models; one per iteration). This is computationally demaning and requires some time. To make sure that this process does not have to run every time the R script is executed, data from parameter optimation process is stored in *.csv* files (eg. *out/decision-tree-eval.csv*) and later loaded into R rather than computed from scratch again. This is the way the script behaves only if the value of *params.readDataFromFiles* is set to *TRUE*.

Task 2a – Decision Tree

Decision tree is a very unstable machine learning algorithm. This means that small differences in training data will result in significant differences in produced models. Therefore high variance is expected when performing the cross-validation.

Only one parameter was tuned: *cp* – the complexity parameter. *cp* defines the minimal decrease in impurity in order for a split to be performed. 50 values in range *0.0001-0.0099* were tested, with step *0.0002*. Each model, produced by changing the complexity parameter, was evaluated by computing the value of $AUC_{0.2}$.

This means 10 values for each value of *cp*. The mean and confidence intervals can be plotted into a line chart with error bars:

Performance of decision tree model for different values of cp

As was predicted, the error bars on the mean of $AUC_{0.2}$ are very large, making it diffucult to select the best value of *cp* which would be statistically significant.

Each value of *cp* can be compared to every other value of *cp* using paired t-test. This way it is possible to calculate how many worse (statistically provable) values of *cp* there are for each value of *cp*. This can be plotted using into a line chart:

Performance of decision tree model for different values of cp

Using this chart it is clear that the best value of *cp* is going to be either 0.0025 or 0.0027:

cp	$AUC_{0.2}$ mean	$AUC_{0.2}$ standard deviation	$AUC_{0.2}$ confidence interval
0.0025	0.0694	0.0121	(0.0607; 0.0780)
0.0027	0.0694	0.0121	(0.0607; 0.0780)

The best value of *cp* is then going to be **0.0025**, because it should produce slightly less overfitted model.

Task 2b – Random Forest

Random Forest is an ensamble machine learning method which uses bagging to create multiple decision trees and then voting to get the final prediction.

In each node only a subset of features is considered as the split candidates. This results in the more important features appearing more often than the less important features. The number of features is controlled by the parameter *mtry*. It is also possible to control the number of trees built using the parameter *ntree*. These two parameters will be tuned.

ntree

The more trees you grow, the better the model; but there is an upper limit to the performance of a Random Forest. The goal when tuning the *ntree* parameter is to find such value of *ntree* that produces a model which can no longer be improved by increasing the number of trees. We will call this value the *Critical Point*.

Random Trees can never overfit the data, thus finding the optimal value of the *ntree* parameter is more about saving computation time than about classifier performance. Also, it is important to find the minimal value of *ntree* before tuning the value of *mtry*.

The tuning began by testing the values given by *round(seq(1, 30) ** 1.6)* to see if we can find the *Critical Point*. The *mtry* parameter was set to default (9). Here is the chart of $AUC_{0.2}$ means and confidence intervals from 10-fold cross-validation:

Performance of random forest model for different values of ntree
x = round(seq(1, 30) ** 1.6)

From this chart it looks like the *Critical Point* might be somewhere between 50 and 70. To analyze this even further, values given by *round(seq(1, 30) ** 1.25)* were tested. The *mtry* parameter was again set to default (9):

Performance of random forest model for different values of ntree
x = round(seq(1, 30) ** 1.25)

From this chart it is even more clear that it is safe to assume that any value of *ntree >**69*** is going to yield good results. The final step was to check if, by any chance, obscurely large values of the *ntree*

parameter produced a better model. For this purpuse values 500, 1000 and 2000 were tested. Here are the results:

ntree	$AUC_{0.2}$ mean	$AUC_{0.2}$ standard deviation	$AUC_{0.2}$ confidence interval
500	0.0583	0.0123	(0.0495; 0.0670)
1000	0.0577	0.0129	(0.0485; 0.0669)
2000	0.0565	0.0130	(0.0472; 0.0658)

As can be seen, the values of $AUC_{0.2}$ have not improved.

mtry

Now, that it is clear that at least 70 trees (>69) in a Random Forest should be enough, we can start tuning the *mtry* parameter. Lower values of *mtry* require more trees in a forest to be effective, so we will work with n*tree=150*.

Values in range {2, 3, 4, 5, … , 26, 27, 28} have been tested to find the best value:

Performance of random forest model for different values of mtry
x = seq(3, 28)

After comparing the performance of each possible value of the *mtry* parameter with the performance of every other value using pairder t-test, we can calculate how many "worse" values there are for each value. The results can be plotted using a line chart:

Performance of random forest model for different values of mtry
x = seq(3, 28)

As can be see from the two charts, the performance improves with higher values of *mtry*. Although values >=15 are impossible to compary decisively. Thus the optimal value of the *mtry* parameter is *>14*.

Task 2c – Regularized Logistic Regression

Task 2d – Model comparison

The parameters of the three models were tuned to find the best values. Then the best values were used to produce three candidate models. The best model of the three candidate models will be selected as the final model, which will then be used for final prediction.
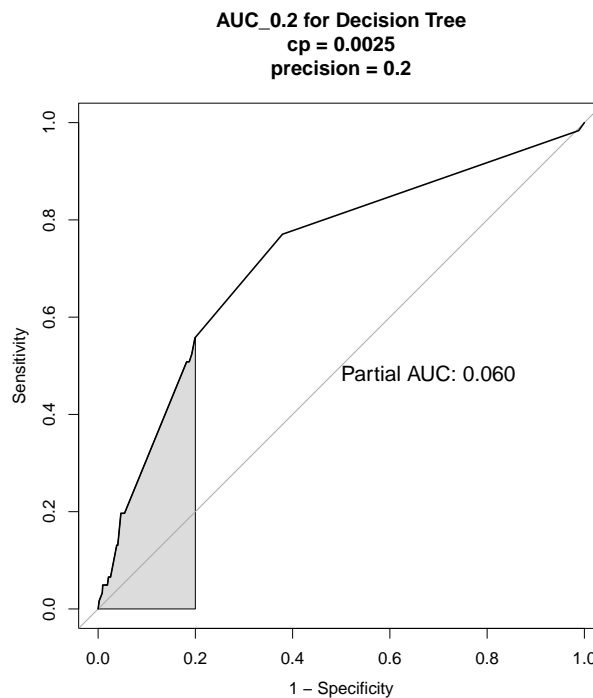
Decision Tree

The best parameters for the Decision Tree model are as follows:

cp	0.0025

The tree that has been build using the value of *cp=0.0025* looks like this:
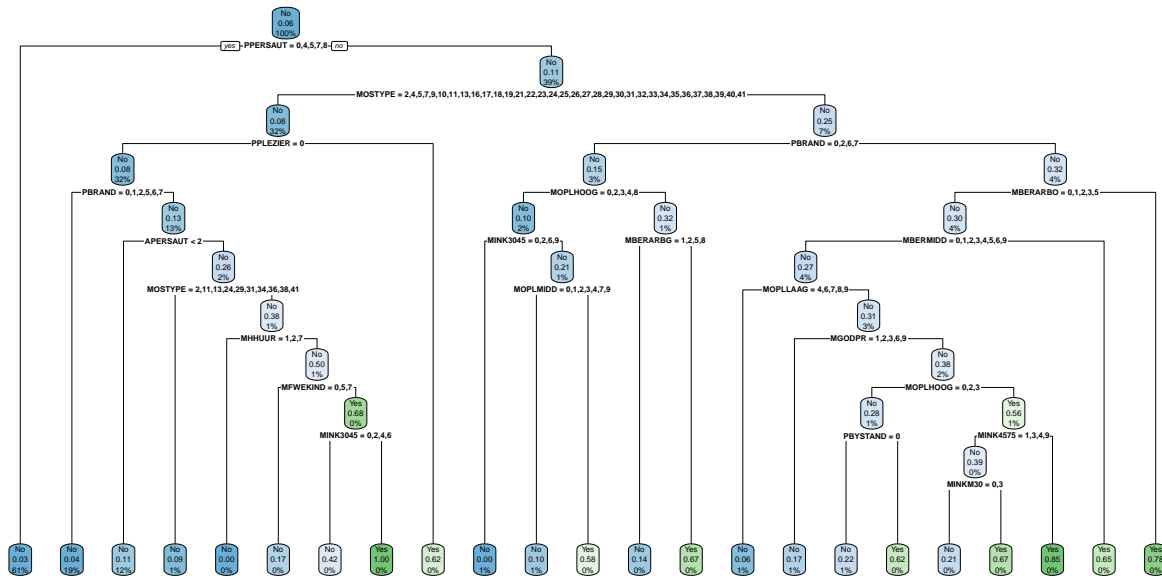
It has been evaluated using an ROC curve. Also, precision when predictng 100 positive examples from 1000 examples in the test data set was calculated:

AUC_0.2 for Decision Tree
cp = 0.0025
precision = 0.2

Partial AUC: 0.060

The Decision Tree model has actually shown to be pretty robus for this classification task.

The five most important attributes according to the Decision Tree model are:

Feature	Importance
MOSTYPE	23.39
PPERSAUT	15.97
APERSAUT	15.08
MFWEKIND	13.50
MFGEKIND	13.04

The topic of variable importance will be later delved into in **Task 3** in more detail.
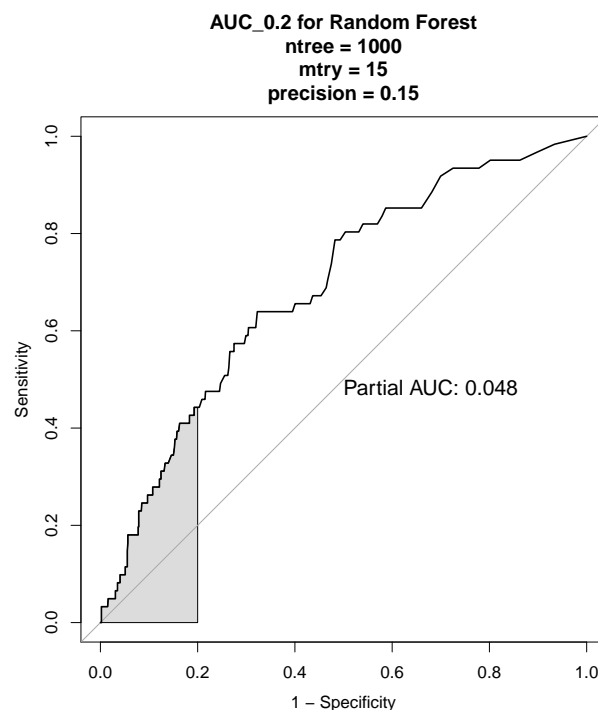
Random Forest

The best parameters for the Random Forest model are as follows:

ntree	>69
mtry	>14

We will use the value of *ntree=1000*, because it cannot hurt us and it will make it easier to explore the feature importance measured by the Random Forest.

The value of *mtry* will be set to *mtry=15*. Lower number of split candidate attributes in each node will make the computation faster.

The Random Forest model with optimal parameters has been evaluated using an ROC curve. Also, precision when predictng 100 positive examples from 1000 examples in the test data set was calculated:

The Random Forest model has suprisingly shown worse results than a single Decision Tree. The reason for this might be that it is really difficult to identify the positive examples and creating an ensamble of very weak classifiers is not going to improve the predictor performance.

The five most important attributes according to the Random Forest model are:

Feature	Importance
MOSTYPE	36.46
PBRAND	22.15
PPERSAUT	16.79
APERSAUT	14.46
MFWEKIND	11.81

The topic of variable importance will be later delved into in **Task 3** in more detail.

Regularized Logistic Regression

Bla Bla Bla

Comparison and final choice

Bla Bla Bla

Task 2e – The best model

Bla Bla Bla

Task 3 – Model interpretation and feature selection

Task 4 – Final prediction on the blind test set

Final prediction on the blind test set was done using the best model with the best parameters. The goal is to be as precise as possible, ie. identify as many true positives as possible from 100 examples marked as positive.